

04_Slices

Slice 0: Bootstrap & Contracts

- Goal: Turn the plan/ADRs into importable code contracts and a runnable skeleton so every later slice builds on stable interfaces.
- Scope
 - Repo scaffold and packaging (one top-level Python package, e.g., bt/).
 - CLI stub with subcommands: bt backtest , bt shadow , bt paper , bt live .
 - Ports/contracts as code (ABC/Protocol):
 - Clock , EventBus (in-process), SnapshotStore , RunManifestStore
 - MarketDataProvider (history + live bars), OrderRouter (paper, real),
 - RiskEngine (pre-trade), PortfolioStore (positions, cash, P&L),
 - Logger/Telemetry (structured logs + metrics),
 - ConfigProvider and SecretsProvider .
 - CI: unit tests + type checks + lint; build and run on push.
 - ADR links/docstrings on each port; versioned with the repo.
- Out of scope:
 - Concrete exchange adapters, full paper/live implementations, non-trivial metrics/alerts.
- Repo structure

```
backtester/
  backtester/core/          # event bus, app services
  backtester/ports/         # interfaces: broker, datafeed, storage, clock
  backtester/adapters/     # paper broker, parquet store, REST polling feed
  backtester/strategy/     # strategy template(s)
  backtester/cli/          # CLI entrypoints
  backtester/sim/          # fill models, slippage, fees
  backtester/risk/
  backtester/stores/       # snapshots, portfolio, parquet
  backtester/util/
tests/
  tests/contract/        # adapter contract tests
  tests/e2e/
sample_data/
ci/
docs/
adr/
scripts/
tickets/
```

- Done when:
 - backtester --help shows the four subcommands; each runs a no-op pipeline.
 - All ports compile and are importable; basic unit test per port passes.

- CI is green: tests + types + lint.
- A sample backtest --noop run writes a `run_manifest.json` and structured logs.

Slice 1: Config & Secrets

- Goal: Make runs reproducible, safe, and auditable via typed config + layered overrides and secret handling.
- **Scope**
 - Typed Config schema (e.g., dataclass) with sections: data paths, symbols, timeframes, costs, risk limits, runtime SLOs, live endpoints.
 - Precedence and merge order: **defaults < file (YAML/JSON) < env vars < CLI flags**.
 - `SecretsProvider`: never logs secrets; supports env vars and one local secure backend (dev).
 - `RunManifestStore`: persist config snapshot, Git SHA, start/end times, and environment (py version, OS).
 - CLI: `--config`, `--set key=value`, and `--env-prefix` support.
- **Out of scope**
 - Cloud secret managers, multi-env promotion flows.
- **Done when**
 - Any run emits a manifest including **full effective config hash** and Git SHA.
 - Missing/invalid config fails fast with a clear error.
 - Secrets never appear in logs or manifests.
 - Re-running the same config on the same data produces the same manifest hash.

Slice 2: Backtest core

- **Goal:** Deterministic backtest loop over **1-minute OHLCV** for one symbol with reliable accounting and auditable outputs. CSV trades; Parquet storage.
- **Scope**
 - **Clock & loop:** deterministic simulated clock stepping bar-by-bar (UTC).
 - **Data ingestion (history):** read Parquet bars; schema: `ts_utc, symbol, open, high, low, close, volume[, ntrades], source`.
 - QC: strict monotonic timestamps, no overlaps; configurable gap handling (`--strict` fail vs `--fill forward/none`).
 - **Event pipeline:** `bar → strategy.on_bar → order intents → risk (minimal) → sim matcher → fills → accounting`.
 - **Matcher (bar-based):**
 - Market fills at bar close ± slippage.
 - Limit fills on cross using side-consistent intra-bar price (min/max).
 - Aggregated fill record per order (MVP).
 - **Accounting:** FIFO lots; realized/unrealized P&L; equity curve; fees/commissions applied post-fill.
 - **Outputs:** `trades.csv`, `orders.csv`, `equity.csv`, `runs/<id>/logs.jsonl`; Parquet bar cache for reproducibility.

- **CLI:** `bt backtest --config path.yml [--from 2024-01-01 --to 2024-03-01] [--strict]`.
- **Perf guardrails:** single-symbol backtest \geq 1 month completes in < 5s on laptop-class HW; memory bounded.
- **Out of scope**
 - Multi-symbol, tick-level, partial-fills fidelity, funding/fees beyond cash spot.
- **Done when**
 - Running the same config/data twice yields **identical file hashes** for `trades.csv` and `equity.csv`.
 - Gap/duplicate bars are **logged** and either **refused** (strict) or **tolerated** (fill mode) per config; choice is visible in the manifest.
 - All outputs are produced; a manifest captures Git SHA, effective config hash, runtime, and data sources.
 - Unit tests cover: bar ordering/QC, limit cross logic, P&L math (FIFO), and determinism (seeded).

Slice 3: Strategy plug + fixed slippage and fees + Minimal Risk in Sim

- **Goal:** Enable strategy development with realistic fills/costs and early pre-trade guardrails—**the same path** later used by paper/live.
- **Scope**
- Strategy API:
 - `on_start(ctx)`, `on_bar(ctx, symbol, bar) -> list[OrderIntent]`, `on_fill(ctx, fill)`, `on_stop(ctx)`.
 - Access to `ctx.state` (`positions`, `cash`), `ctx.now_utc`, and `ctx.submit(order)`.
- Fill/Execution policy (sim):
 - Market: fill at bar close \pm slippage.
 - Limit: fill if price crosses; price = min/max within bar consistent with side.
 - Partial fills aggregated into a single fill record in sim (MVP simplification).
- Cost model: pluggable `CostModel` with commissions/fees, spread, slippage (bps or function), min notional.
- **Minimal Risk (enabled in sim now; reused later):**
 - Allowed symbols whitelist.
 - Per-order notional cap.
 - Max open positions count.
 - Daily loss limit (hard stop \rightarrow block new orders).
 - All risk blocks are logged with reason codes.

Out of scope

- Advanced exposure buckets, leverage/margin logic, borrow/short.

Done when

- Deterministic backtest: given same data+config, orders and P&L are identical run-to-run.
- Costs are included and can be toggled/parameterized via config.
- Risk violations block orders in sim and are visible in logs with reason codes.

- Unit tests cover: limit-cross logic, slippage application, and each minimal risk rule.

Slice 4 Paper trading parity

- **Goal:** Run on **live 1-min bars** with an in-process paper broker using the same order path as sim, to validate parity before any real exchange.
- **Scope**
 - **LiveDataProvider (REST polling)**: UTC timestamps, retries/backoff, gap detection, lag metric.
 - **PaperOrderRouter**: same state machine as live (new/submitted/filled/canceled/rejected); idempotent `client_order_id`; restart-safe.
 - **Fill policy**: identical to sim (market @ close \pm slippage; limit on valid cross).
 - **State**: PortfolioStore persisted; restart restores positions/cash/open orders.
 - **Controls**: global **halt/flatten**; dry-run mode logs would-be actions only.
 - **Telemetry**: `bars_lag`, `strategy_tick_latency`, `orders_submitted`, `orders_blocked_risk`.
- **Out of scope**:
 - Real exchange connectivity, advanced alerting dashboards
- **Done when**
 - Same-day **paper vs recorded sim** P&L difference **median ≤ 5 bps**.
 - **No missed bars** over a continuous **7-day** paper run (gaps detected & logged).
 - Idempotent order IDs prevent duplicates across restarts.
 - **Latency SLOs** met end-to-end: **p95 ≤ 500 ms, p99 ≤ 900 ms** from bar arrival to decision.

Slice 5 Shadow Mode

- **Purpose:** Soak-test the live runtime safely: compute signals and risk-check orders, but **do not submit** anywhere. Capture logs/metrics to validate stability and SLOs.
- **Scope**
 - `NullOrderRouter` that records proposed orders (with reasons if blocked) but never submits.
 - Full live pipeline: data \rightarrow strategy \rightarrow risk \rightarrow (null) router \rightarrow logs/metrics/snapshots.
 - Health and observability:
 - Metrics: `bars_lag`, `strategy_tick_latency`, `order_proposed_rate`, `risk_block_rate`, `error_rate`.
 - Alerts/warnings (log-level is fine for MVP): bar gap detected, error rate spike, daily loss threshold hit (from shadow P&L).
 - Artifacts: daily run manifests and compressed logs for replay.
- **Out of scope:** Exchange auth, order placement, reconciliation with real positions.
- **Done when**
 - **7-day** continuous shadow run completes without missed bars (gaps reported \leq configured tolerance).
 - Latency SLOs met: **p95 ≤ 500 ms, p99 ≤ 900 ms** decision latency.
 - Metrics show non-zero proposed orders and risk blocks (i.e., strategy is “doing work”).

- Crash/restart resumes from the last snapshot with consistent internal state.

Slice 6 — Risk rails & safety

- Goal: Enforce per-order and portfolio guardrails pre-trade; provide an operator **kill-switch** to halt and flatten safely.
- **Scope**
 - **Rules (pre-trade):**
 - Allowed symbols whitelist.
 - **Per-order notional cap** (quote currency).
 - **Max open positions** count.
 - **Daily realized loss limit** → hard stop (blocks new orders).
 - **Evaluation:** deterministic order; first failing rule stops evaluation; attach **stable reason code**.
 - **Operator controls:**
 - `bt halt` (block new orders immediately).
 - `bt flatten` (submit closes for all open positions, then halt).
 - Manual resume requires explicit flag.
 - **Observability:** per-rule counters, last block reason, rolling daily loss.
 - **Configuration:** thresholds in typed config; rules can be toggled per env (sim/paper/live).
- **Out of scope**
 - Leverage/margin, shorting/borrow, sector/bucket exposures (non-crypto).
- **Done when**
 - Any order breaching **notional**, **max positions**, or **daily loss** is blocked; an audit log entry with **reason code** is produced.
 - **Kill-switch** halts within **≤ 1s**; **flatten** closes positions then sets halt state; resume requires manual flag.
 - Shadow/paper runs show non-zero **risk_block_rate** under forced breach scenarios.

Slice 7 Binance REST live adapter (use Binance wrapper)

- **Goal:** Submit **tiny real orders** with the same semantics as paper; robust to restarts, rate limits, and transient failures. Use REST first (polling), then add WebSocket later (non-MVP).
- **Scope**
 - **Auth & signing:** API key/secret via `SecretsProvider`; HMAC signing; clock skew guard (time sync endpoint).
 - **Order placement:** `clientOrderId` idempotency; map venue filters (`LOT_SIZE`, `MIN_NOTIONAL`, `PRICE_FILTER`) at submit-time; reject early on violation.
 - **State & reconciliation:**
 - Poll open orders & account snapshot; reconcile against local state on startup and periodically.
 - Aggregate venue partial fills into **single logical fill** upstream (MVP policy).
 - **Error handling:** map HTTP/JSON errors to stable codes; classify **retryable vs terminal**; exponential backoff; respect rate-limit headers.

- **Safety**: tiny size orders; configurable symbol allow-list; auto-halt on repeated failures.
- **CLI**: `bt live --adapter binance --tiny --config path.yml`.
- **Out of scope**: WebSocket streaming, margin/derivatives, advanced order types beyond market/limit GTC.
- **Done when**
 - End-to-end tiny orders **fill** using the same state machine as paper.
 - **Startup reconciliation** yields correct positions; no duplicate live orders after process crash.
 - **Rate-limit handling** verified (throttled path exercised) with backoff; no hot loop.
 - **Venue filter violations** are blocked pre-submit with clear reason.
 - p95 order **round-trip ≤ 500 ms**, p99 ≤ 900 ms (from submit to ACK/terminal event) under light load; no missed live bars for a week.

Slice 8 Observability & recovery

- **Goal**: Metrics, structured logs, **snapshots/checkpoints**, and **crash-replay**; minimal alerts for bar gaps, order errors, and daily loss halts.
- **Scope**
 - **Structured logging**: JSONL with run id, ts_utc, component, level, msg, context (order id, symbol).
 - **Metrics**: `bars_lag`, `strategy_tick_latency`, `order_rtt`, `error_rate`, `risk_block_rate`, `pnl_intraday`.
 - **Snapshots**: state snapshot on position change and at fixed cadence; include open orders and cash.
 - **WAL**: append-only order/decision log enabling idempotent replay.
 - **Replay**: `bt replay --from snapshot --wal path` reconstructs state and verifies idempotency.
 - **Alerts** (log-level for MVP):
 - **Bar gap**: consecutive bar gap > tolerance.
 - **Order error rate**: rolling window above threshold.
 - **Daily loss halt**: triggered by RiskEngine.
- **Out of scope**
 - External monitoring stacks (Grafana/Prom), pager duty, rich dashboards.
- **Done when**
 - **Crash mid-submit** test: recovery results in correct state and **zero duplicate live orders**.
 - Metrics above are emitted; three alerts fire under induced conditions and are visible in logs.
 - Daily CSV/Parquet rollups generated; minimal summary printed at end of each run.

Contract Tests (Gating Before Real Orders)

- **Purpose**: Prove every adapter and store honors its port contract. These are **automated**, **repeatable** checks that gate enabling real trading.

- **Scope**

- **MarketDataProvider contract tests**

- History is chronological, **UTC**, complete for the requested window (or gaps are explicitly surfaced).

- Live bars' timestamps are strictly increasing; lag metric reported.

- **OrderRouter contract tests** (run against Paper and any real adapter sandbox/demo):

- `submit` returns ACK with idempotent behavior on duplicate `client_order_id`.

- `cancel` transitions correctly; cancel after fill is a no-op with stable status.

- Rate-limit and transient errors map to **retryable** codes with backoff.

- Partial fills (if any from venue) are **aggregated** into a single final fill event presented upstream (MVP policy).

- **RiskEngine contract tests**

- Each rule blocks when it should and emits a stable reason code.

- **SnapshotStore / PortfolioStore contract tests**

- Save/restore yields byte-for-byte equivalent trading state (positions, cash, open orders).

- **Clock/EventBus**

- Deterministic scheduling in sim; subscribers receive events in publish order (single-threaded MVP).

Tooling & Run

- CLI: `bt contract-tests [--adapter paper|binance-sandbox|...]`

- Runs locally in CI; any failure blocks merging/enabling real orders.

Done when

- All contract tests pass for: MarketDataProvider (history, live), OrderRouter (paper), RiskEngine, SnapshotStore, PortfolioStore, Clock/EventBus .
- A single summary report is produced (JUnit/JSON) and stored as an artifact in CI.
- Before enabling a real adapter, it must pass the same suite against its sandbox/demo environment.